

# Plan 9

A New Paradigm for Operating Systems



| term% fortune  
| Your fantasy will come true.





# A Brief History

- Began development in late 80s
  - Development led by a star-studded cast
    - Rob Pike, Brian Kernighan, Ken Thompson, Bjarne Stroustrup
  - Released to public in 1995
  - Fine-tuned for research and networking – not user-friendly
  - Unix: “Everything is a file”
    - Minimalism, portability, purity
  - Laid to rest Jan 10, 2015
  - Carried on by 9front, Inferno, 9legacy, etc.

```
pu% apl
d+2 3p?^6
d
) 5 1
) 4 6
sort<=c\$A]-
sort d
) 4 6
) 5 1
sort>=l-d
) 3 5
) 4 6
+d#
) 9 7
+/d
) 12
t2 3
```


1	1	1	2	1	3
2	1	2	2	2	3





# Idiosyncrasies

- Independent C library
  - No GNU, and no GCC! A dialect of C.
- Custom, but portable, C compilers
  - Unique compiler for each architecture
  - 8c for x86, 6c for AMD, 5c for ARM, 0c for MIPS
- 9P network protocol
  - Used by Microsoft for WSL since 2019!
- Independent windowing system (rio)
  - Doesn't use X Window System
- A distributed operating system
  - Total data communication transparency
  - Encourages cluster computing
- Not Bash, not PowerShell, but rc!
  - Some commands are the same, many are different
  - No symbolic links, and no sudo (it's unsafe!)



```
term% lc
bin    hello.c lib      tmp
term% 8c hello.c
term% 8l hello.8
term% lc
8.out  hello.8 lib
bin    hello.c tmp
term% 8.out
hello world
term%
```



# Synchronization via channels

- Threads communicate via channels
- Mutexes, condition variables, and locks are not necessary
  - Although they are provided
- **send()**: sends value through channel, blocks
- **recv()**: receives value through channel, blocks
- **chancreate()/chanclose()** enforce synchronization – no primitives needed
- Channels are powerful enough to implement pipes, sockets, and semaphores

```
term% lc
acme.dump    lib      readme.acme tmp
bin          proj     readme.rio

term% cd proj
term% ls
8.out
proj.8
proj.c

term% 8c proj.c
term% 8l proj.8
term% ./8.out
started lockstep up to 20 iterations
initiator: 1
responder: 2
initiator: 3
responder: 4
initiator: 5
responder: 6
initiator: 7
responder: 8
initiator: 9
responder: 10
initiator: 11
responder: 12
initiator: 13
responder: 14
initiator: 15
responder: 16
initiator: 17
responder: 18
initiator: 19
responder: 20
initiator detected max count (20), closing channel
responder: channel closed: exiting.

term%
```

```
Newcol Kill Putall Dump Exit
New Cut Paste Snarf Sort Zerox Delcol
proj.c Del Snarf Undo | Look!
// demonstrates channels

#include <u.h>
#include <libc.h>
#include <thread.h>

#define MAX_VALUE 20

Channel* initiator_channel;
Channel* responder_channel;

void
initiatorthread(void*)
{
    int control;

    for(;;) {
        recv(initiator_channel, &control);

        if (control >= MAX_VALUE) {
            print("initiator detected max count (%d), closing channel\n", control);
            chanclose(responder_channel);
            threadexits(nil);
        }

        control++;
        print("initiator: %d\n", control);
        send(responder_channel, &control);
    }
}

void
responderthread(void*)
{
    int control;

    for(;;) {
        if(recv(responder_channel, &control) < 0) {
```



# Concurrency in Plan 9

- Same capabilities as POSIX threads, but with a different library
- The iconic producer-consumer problem
- Again, we use channels to avoid race conditions
- Channels are “typed”, unlike Linux
  - Recall: `send()` vs. `sendul()`, and `recv()` vs. `recvul()`
  - Untyped (byte stream, pointers) vs. typed (direct, pass by value)
- Unlike Linux, no thread attributes, no joining
- Writers and readers automatically synchronize thanks to rendezvous communication

```
term% lc
acme.dump    lib
bin          proj
term% cd proj
term% 8c proj.c
term% 8l proj.8
term% ./8.out
p 1
p 2          c 1
p 3          c 2
p 4          c 3
p 5          c 4
p 6          c 5
p 7          c 6
p 8          c 7
p 9          c 8
p 10         c 9
p 11         c 10
p 12         c 11
p 13         c 12
p 14
```

```
Newcol Kill Putall Dump Exit]
New Cut Paste Snarf Sort Zerox Delcol]
+Errors Del Snarf | Look
[Del Snarf Undo | Look
#include <u.h>
#include <libc.h>
#include <thread.h>

#define MAX_VALUE 20

Channel* initiator_channel;
Channel* responder_channel;

void
initiatorthread(void*)
{
    int control;

    for(;;) {
        recv(initiator_channel, &control);

        if (control >= MAX_VALUE) {
            print("initiator detected max count (%d),\n"
                  "closing channel\n", control);
            chanclose(responder_channel);
            threadexits(nil);
        }

        control++;
        print("initiator: %d\n", control);
        send(responder_channel, &control);
    }
}

/usr/glenda/proj/proj.c Del Snarf Get | Look]
```

```
Newcol Kill Putall Dump Exit]
New Cut Paste Snarf Sort Zerox Delcol]
/usr/glenda/ Del Snarf Get | Look
acme.dump    lib/      readme.rio
bin/          readme.acme tmp/
[Del Snarf Undo | Look
void
responderthread(void*)
{
    int control;

    for(;;) {
        if(recv(responder_channel, &control) < 0) {
            print("\t responder: channel closed: exitin
g.\n");
            threadexits(nil);
        }

        control++;
        print("\t responder: %d\n", control);
        send(initiator_channel, &control);
    }
}

void
threadmain(int, char*[])
{
    int start;

    initiator_channel = chancreate(sizeof(int), 0);
    responder_channel = chancreate(sizeof(int), 0);

    threadcreate(initiatorthread, nil, 8*1024);
    threadcreate(responderthread, nil, 8*1024);

    start = 0;
    print("started lockstep up to %d iterations\n",
          MAX_VALUE);
    send(initiator_channel, &start);

    threadexits(nil);
}
```



# System Calls

- Note the presence of a kernel-side **note system**
  - *Pun not intended*
- Single words of data may be exchanged via **rendezvous**
- Plan 9: 38 syscalls; Linux: Over 400
- Note: Plan 9 has no kernel threads!

Memory management	
brk_	allocate memory
segattach	attach to or create a segment
segdetach	detach from a segment
segfree	free physical memory
segbrk	change segment length
segflush	flush cache

Notes	
alarm	set the alarm timer
notify	set the note handler
noted	continue after note

Files	
open	open an existing file
create	create a new file or open an existing file for writing
pread	read from an open file
pwrite	write to an open file
chdir	change current directory
seek	change the current position in an open file
close	close a file descriptor
dup	duplicate a file descriptor
fd2path	retrieve file name
stat	read file metadata
fstat	read open file metadata
wstat	write file metadata
fwstat	write open file metadata
remove	delete a file

Namespace management	
mount	mount a 9P connection
bind	bind a file or directory
umount	unmount or remove a bind
9P connections	
fversion	initialize a 9P connection
fauth	initiate authentication
Process management	
rfork	change process attributes or create a new process
exec	replace the current process
exits	terminate the current process
errstr	exchange error string
sleep	sleep a given amount of time
Synchronization	
await	wait for a child process to terminate
pipe	create a pipe
rendezvous	exchange a word of data
semacquire	acquire a semaphore
semrelease	release a semaphore



# Significance and Impact

- Once-native Plan 9 development tools, like `rc`, `acme`, and `plumbing` have been ported to Linux, Mac, and Windows
- The 9P filesystem is made use of today by Windows' WSL
- Served as a catalyst for the software minimalism movement
- A pioneering system in grid computing
- Served as the backbone for years of innovative research at Bell Labs
- Ran by Coraid hardware (data storage)
- Influenced Google's Go-lang
- Inspired Inferno, 9front, and other OSes

